Imperial College
London

IMPERIAL COLLEGE LONDON

DEPARTMENT OF COMPUTING

Chameleon Text: Exploring ways to increase variety in artificial data

Author:
Thien P. Nguyen

Supervisor:
Julia Ive

Abstract

Variational autoencoders (VAEs) have shown a lot of promise in other machine learning fields but are relatively unexplored in comparison to their potential applications as a language model. We explore the potentials of VAEs within the NLP domain for the purposes of data represention of various datasets, representing a variety of sequence based problems. In solving this problem, we encounter issues from both the sequentiality of recurrent models and also variational autoencoders, and discuss approaches that contemporary models have used to circumvent them. Emperical experiments conducted on the Amazon Reviews, OpenSubtitle and Penn-Tree Bank datasets show how the proposed models encompasses more information within the latent parameters as opposed to earlier models.

Contents

Chapter 1

Introduction

Artificial datasets for NLP purposes is often left unexplored as compared to the field of machine learning in general. For instance, generative models have advanced greatly in the visual domain. That being said, it can be considered to have a demand (and similarly so for any type of data) due to the fact that quality, labelled datasets have limited availability, or that resources towards making said quality datasets are not viable.

Recently, variational bayesian models have shown potential from both theoretical and practical perspectives (Kingma and Welling (2013)). We explore the capabilities of one of the more recent advances in neural sequence modelling by Du et al. (2018), called the VAD. The nature of this paper is exploratory, and not necessarily to introduce novel contributions to the field. In this paper, we review historical implementations, the theoretical and emperical properties of the VAD, and evaluate their applicability for sequence generation within the domain of proposed datasets by He and McAuley (2016), Lison and Tiedemann (2016). We conclude with future steps towards reaching a solution to the problem domain.

Chapter 2

Literature Survey

To provide context into the literature survey, we imagine a scenario where a client requests the use of our dataset. Exposing the client access to the data would be inappropiate for multiple reasons, ranging from the exposure to the privacy of the users within the dataset and their idiosyncracies, to the impracticalities of transferring large datasets. The objective is to provide some alternative dataset to ours such that value can be deduced from the data, but the privacy of the users in our original dataset is maintained.

This can be accomplished by creating a language model that would encompass the lexical variety of our original dataset. This language model would train on our original dataset, and would produce data that is semantically and lexically similar to our original data, but diverges enough such that it could potentially be seen as an entirely new and independent dataset. This new "dataset" can be provided to other clients as they would be able to generate responses from the model, which from there they can deduce higher level properties of the dataset, without exposing the idiosyncracies that created the original properties of the dataset.

The premise of this literature survey is to explore the core components necessary to construct our solution, and to also discuss approaches to the problem.

2.1   Text Generation

Text generation is a a type of language modeling problem, which in itself, is one of the core natural language processing problems. It is typically considered challenging as samples are discrete and results are non-differentiable (Kovalenko (2017a); Kovalenko (2017b)), as opposed to other types of data mediums.

Traditionally, data itself could be generated via the use of data augmentation techniques. This is primarily the case for images; For instance, images can be indiscriminately flipped, cropped and transformed. These operations cannot be extended to sequences of text as words are temporally and contexturally bound, such that the omission to certain words in a sequence could dramatically change the semantics of the sequence. Even if these methods could be successfully applied, the resulting data would be relatively rudimentary in terms of their throughput, lexical diversity, and coherence.

With the advent of neural language models (explained in detail later on), text generation became more viable, and can now be seen in a variety of applications, ranging from machine translation (Sutskever et al. (2014)), to email response generation (Kannan et al. (2016)), to document summarisation (Nallapati et al. (2016)).

2.2 Language Modelling

Language modelling is the task of predicting a word w_i in a text w given some sequence of previous words $(w_1, w_2, ..., w_{i-1})$. More formally, Dyer (2017) describes an unconditional language model as assigning a probability to a sequence of words, $w = (w_1, w_2, ..., w_{i-1})$. This probability can be decomposed using the chain rule:

$$p(w) = p(w_1) \times p(w_2|w_1) \times p(w_3|w_1, w_2) \times ... \times p(w_i|w_1, w_2, ..., w_{i-1}) \tag{2.1}$$

$$p(w) = \prod_{t=1}^{|w|} p(w_t|w_1, ..., w_{t-1}) \tag{2.2}$$

Traditionally, assigning words to probabilities may conflate syntactically spurious sentences (e.g. it may be the case that predicting two verbs together is improbable but still technically possible) but it remains to be a useful method for representing texts. For the purposes of this paper, we concentrate on conditional language modelling. A conditional language model assigns probabilities to sequences of words, $w = (w_1, w_2, ..., w_{i-1})$, given some conditioning variable, x.

$$p(w|x) = \prod_{t=1}^{|w|} p(w_t|x, w_1, ..., w_{t-1}) \tag{2.3}$$

There exists a variety of language models; we start with the n-gram, argued as being the most fundamental (Le (2018)). An n-gram is a chunk of n consecutive words in a sequential order. For instance, given the sentence "the quick brown fox . . . ", the respective n-grams are:

- unigrams: "the", "quick", "brown", "fox"

- bigrams: "the quick", "quick brown", "brown fox"

- trigrams: "the quick brown", "quick brown fox"

- 4-grams: "the quick brown fox"

The intuition of n-grams was that statistical inference can be applied on the frequency and distribution of of the n-grams, which could be used to predict the next word. That being said, one of the caveats of n-grams is that sparsity is not captured, and is not necessarily viable to. For example, increasing the sparse relationships between words can be captured by increasing the n-gram size, but this increases

the odds of capturing an n-gram of zero probability. There exists mechanisms to bypass this issue (such as Smoothing, Backoff and Interpolation Jurafsky and Martin (2019)), but using n-grams still remain infeasible as using such mechanisms does not negate the potential data requirements for increasingly larger n-grams.

Later on, modern language models began to revolve around the use of neural networks (Bengio et al. (2001)), which itself started off with a multi-layer perceptron (MLP) that formed the premise of word prediction. The use of neural networks in language modelling is often called Neural Language Modelling, of NLM for short.

Neural Networks are non-linear and computational metaphors of the brain that model arbitary relationships between input and output vectors. A textbook architecture of a neural network revolves around the multi-layer perceptron. Note that the input and output vectors are of a fixed dimension. Neural Networks evaluate an input using forward propagation, to produce an output. Traditionally, neural networks are trained to produce optimal outputs via the use of backpropagation.

2.3 Recurrent Neural Networks

Recurrent neural networks (RNNs) are a class of feed forward neural networks such that the outputs are not necessarily restricted and discrete in dimension (as opposed to the MLP). RNNs operate over some variable-length vector, and produces an output of some other arbitary variable-length vector. This circumvents a problem introduced with using an MLP, where sentences are not typically fixed in length.

The architecture of RNNs make it favourable in NLP related problems as words in sentences are typically conditioned on the previous words. When treated as language models, a standard RNN language model predicts each word of a sentence conditioned on the previous word and an evolving hidden state.


Figure 2.1: Rectangles represent vectors, with red being the input, blue the output, and green representing the state of the RNNs. Arrows represent functions. From left to right: (1) an MLP. (2,3,4,5) are examples of different styles of recurrent neural networks, describing the different types of input and output combinations. (Diagram from Karpathy (2015))

At each time step t, a simple RNN produces a hidden vector h_t, derived from the input vector x_t and the previous state h_{t-1} in the function $h_t = f_w(h_{t-1}, x_t)$. The hidden vector is usually obtained by some affine transformation offset by a bias vector, described in more detail in equation 2.4. The same function and the same parameters are used at each time step t. RNNs, by design, are not computationally

parallelisable. This hidden vector h_t is continously updated as it traverses through the RNN, learning the temporal relationships between the inputs.

$$h_t = tanh(W_{hh}h_{t-1} + W_{xh}x_t)$$
$$y_t = W_{hy}h_t$$

(2.4)

Equations for the RNN cell.

RNN cells can be stacked, making a deeper recurrent network, and can also be bi-directional. Bi-directionality explores the temporal relationships of the sequence in both directions (left to right, and vice versa.) This is made possible by stacking two recurrent cells on top of each other, and having one forward propagate in one direction, and another in the opposing direction.

RNNs are typically trained with backpropagation through time (see equation 2.5) The gradient flow $w \leftarrow w - \alpha(\delta L/\delta w)$ is computed at every timestep, meaning that every path from W to the loss $\mathcal{L}$ needs to be considered.

$$\frac{\delta L}{\delta w} = \sum_{j=0}^{T} \sum_{k=1}^{j} \frac{\delta L_j}{\delta y_j} \frac{\delta y_j}{\delta h_j} \left(\prod_{t=k+1}^{j} \frac{\delta h_t}{\delta h_{t-1}} \right) \frac{\delta h_k}{\delta w}$$

(2.5)

Typically, there are two caveats with training RNNs. One of which involves the sensitivity of the gradients: The product symbol within the loss function in 2.5 is the cause of vanishing and exploding gradients, both of which make it especially difficult to train RNNs. This is compounded by the fact that the limited expressibility of the gradients means that relationships between timesteps earlier in the recurrent computation are harder to train than timesteps further in the sequence. This is described as the long range dependency problem.

Additionally, it became aparrent that it was difficult for RNNs to establish relationships between potentially relevant non-contiguous inputs and outputs, especially in the case for longer sequences, there isn't necessarily a clear indicator in the architecture that would facilitate this feature.

That being said, RNNs remains to be appropiate for language modelling primarily for its capability of encapsulating the temporally contextual relationships within the input sequence. The idea of recurrent language models can be realised in Section 2.6. Our paper revolves around the use of recurrent language models in a generative manner. We describe the generative capabilities further in Section 2.4.1.

2.3.1   LSTMs and GRUs

LSTMs and GRUs are types of recurrent cells, introduced to circumvent the issue of long range dependency problems, and gradient sentivities that RNN cells suffered. LSTM (Long Short Term Memory; Hochreiter and Schmidhuber (1997)) are a type of RNN cell that retains information based on the previous inputs through the introduction of gated architectures. These gated architectures regulate the flow of

$$
\begin{aligned}
i_t &= \sigma(W_i \cdot [h_{t-1}, x_t] + b_i) \\
f_t &= \sigma(W_f \cdot [h_{t-1}, x_t] + b_f) & u_t &= \sigma(W_u \cdot [h_{t-1}, x_t] + b_u) \\
o_t &= \sigma(W_o \cdot [h_{t-1}, x_t] + b_o) & r_t &= \sigma(W_r \cdot [h_{t-1}, x_t] + b_r) \\
g_t &= \sigma(W_g \cdot [h_{t-1}, x_t] + b_g) & c_t &= tanh(W w_{t-1} + U(r_t \odot h_{t-1})) \\
c_t &= f_t \odot c_{t-1} + i_t \odot g_t & h_t &= (1 - u_t) \odot h_{t-1} + u_t \odot c_t \\
h_t &= o_t \odot tanh(c_t)
\end{aligned}
\tag{2.6}
$$

Figure 2.3: Equations for LSTM cells (left) and GRU cells (right).

information coming to and from the cell, which helps to mitigate the issue of gradient sentivity and improves long range dependencies. Gates of the cell are parameterised and therefore learnable. LSTM cells can be swapped in place with the RNN cells. It has shown to have a considerably stronger performance for a variety of tasks compared to RNNs cells.

GRUs (Gated Recurrent Units; Cho et al. (2014)) are functionally similar to LSTMs, but uses less parameters. Within the GRU architecture, a feature to retain the previous weights remain, but there exists an direct path to the input data from the output, allowing a reduction in training time. Chung et al. (2014) suggests that GRUs were found to perform at least equivalently to LSTMs but show slight improvements on smaller datasets. As there are objectively less parameters for the GRU, it should converge faster than LSTMs.

2.4 Autoencoders

Autoencoders (E. Rumelhart et al. (1986)) are a specialised form of neural networks where the model attempts to faithfully recreate the inputs on the output. This is made possible by learning the distribution of the input data. Autoencoders typically have a layer in the model where its dimension is smaller than the input space, therefore representing a dimensionality reduction in the data. Autoencoders can be considered to be a non-linear representation of PCA. Autoencoders within the domain of NLP are popularised through their use in Machine Translation, Word Embeddings, and document clustering.

Autoencoders are composed of two different principal components, an encoder network α and a decoder network β, such that $\alpha : X \rightarrow F$ and $\beta : F \rightarrow X$. Note that F is typically smaller in dimension than X, otherwise the model would be learning the identity vector. Measuring the success of the reconstruction is deduced by a reconstruction loss formula. This reconstruction loss comapres the output of the decoder and compares it against the input of the encoder. The two networks are trained together in a manner that allows them to preserve the input as much as possible in the output.

Although the model learns the distribution of the data, unlike PCA one cannot immediately interpret the principal components of the data. In fact, it is difficult to

understand the structure of the latent variable, as is the case for understanding the layers of neural networks in general. Additionally, as there is no sampling mechanism within this model, this effectively prohibits any generative capabilities.

2.4.1 Variational Autoencoders

Variational Autoencoders (VAEs, Kingma and Welling (2013)) stem away from regular autoencoders. As opposed to learning a fixed latent vector, VAEs force the encoder to learn parameters of a gaussian distribution - a mean μ and a covariance Σ. In more discrete terms, it replaces the encoder from the standard autoencoder with a learned posterior *recognition model*, $q(z|x)$. This component of the autoencoder parameterises an approximate posterior distribution over z, using a neural network conditioned on x. This consequently allows us to take advantage of recent advances in variational inference. A sample is taken from a gaussian using the learned parameters, which is then fed into the decoder. This sampling mechanism allows VAEs to have generative capabilities. For instance, one can entirely avoid the encoder network, and augment the mean and variance vectors to generate outputs from the decoder.
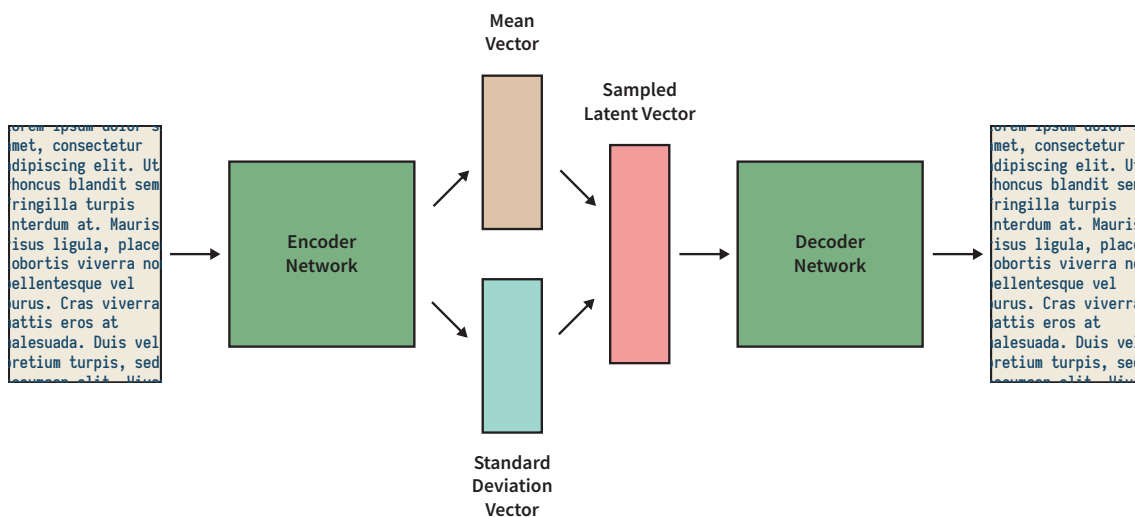

Figure 2.4: An abstracted model architecture for a variational autoencoder, which takes as input some text, and it's predicted output being the same text as the input.

Note that the decoder receives samples from a non-standard normal distribution produced by the encoder. The average of the samples of the different distributions should approximate to a standard normal. This stochasticity allows us to model variability in the results.

If VAEs were trained with a standard autoencoder's reconstruction objective, then it would learn to encode its inputs determinsitically by making variances in the *recognition model*, $q(\overrightarrow{z}|x)$ vanishingly small. (Raiko et al. (2014)). When measured with KL Divergence, creates a obsolete value - This is known as a vanishing KL. (Fu* et al. (2019)) Instead, the loss function for VAEs is composed of two components; a reconstruction loss that involves an expection of the output; and a KL divergence (see

Equation 2.8), which measures the distribution of the posterior distributions against a standard gaussian $\mathcal{N}(0,1)$. This objective provides a tractable lowerbound on the true likelihood of the data, and forces the model to encode some information on the data into the gaussian parameters.

$$\mathcal{L}(\theta, \phi, x, z) = \mathbb{E}_{q_\phi(z|x)}[log\, p_\theta(x|z)] - D_{KL}(q_\phi(z|x)\,||\,p(z)) \leq log(p(x)) \tag{2.7}$$

$$D_{KL}(P||Q) = \sum_{x \subset X} P(x) \cdot log(\frac{P(x)}{Q(x)}) \tag{2.8}$$

Reparameterisation Trick

$$z = \mathcal{N}(\mu, \Sigma) \equiv \mu + \Sigma \cdot \epsilon \tag{2.9}$$

Calculating gradients through backpropagation would be impossible as gradients would have to eventually push through a randomly sampled latent variable. The gaussian reparameterisation trick (Kingma and Welling (2013)) subverts this issue through the rearrangement of the gaussian parameters. Instead of sampling z through a gaussian $\mathcal{N}(\mu, \Sigma)$, z can be created rearranging the definition of the variable; a random sample through a gaussian can be redefined with a summation of μ and some gaussian noise $\epsilon \sim \mathcal{N}(0,1)$ applied to Σ. This rearrangement allows for derivatives to be calculated with respect to the parameters that are needed for learning as derivatives with respect to the relevant components would cancel out ϵ.

2.4.2 Conditional Variational Autoencoders

In a VAE, the decoder cannot produce outputs of a particular condition on demand. CVAEs add to the VAE to include conditioning on another description of the data, a class descriptor y.


Figure 2.5: A model architecture for a CVAE, which includes the label being fed into the encoder and decoder networks.

During training time, a class (represented by some arbitary vector) is fed at the same time to the encoder and decoder. To generate an output that depends on y we feed that datum to the decoder along with a random point in the latent space sampled from a standard normal distribution.

Samples can be generated from the conditional distribution $p(x|y)$. By changing the value of y, we can get corresponding samples $x \sim p(x|y)$. The system no longer relies on the latent space to encode what output is necessary; instead the latent space encodes other information that can distinguish itself based on the differing y values.

In terms of implementation, it happens to be more feasible to concate the class variable to the dataset prior to feeding. This allows the model to retain its characteristics without needing to adjust in order to accomadate the conditioning variables.

2.5 Related Work

2.5.1 Sequence To Sequence

Sequence to Sequence, (seq2seq, Sutskever et al. (2014)) is a type of neural language model that models relationships between sequences. Seq2Seq comprises of two components, encoders and decoders (in a similar fashion to autoencoders), both of which are represented with RNN models (although as explained earlier, can be replaced with LSTM or GRU cells). The data it takes is a pair of sequential data, which for the example of text generation, would be query and response sequences. The query sequence would go through the encoder model. The last hidden state value of the encoder is used as the context variable, which encodes the characteristics of the input sequence. This is used as the first hidden vector value for the decoder, which would then predict the response sequence. For the purposes of this paper, we describe this as the standard recurrent discourse model.

Figure 2.6: An abstracted model of the seq2seq architecture, where the encoder (pink) takes in the input sequence, and the decoder (blue) shows the output sequence. The encoder outputs are effectively ignored.

Seq2seq models have shown to be effective in reponse generation, but is not without it's own set of problems. For instance, the context variable itself is considered to be the bottleneck as it has to encompass the properties of the entire input sequence in a single vector. As we've discussed earlier in Section 2.1, the training mechanism is

subject to vanishing and exploding gradients, and that it becomes increasingly difficult to retain information about earlier parts of the sequence as opposed to recent parts of the sequence in the encoder stages as a consequence. We explain later about approaches to mitigate the issue.

Furthermore, the responses produced from a seq2seq model naturally lacks lexical diversity and richness (Serban et al. (2016), Zhao et al. (2017), Jiang and de Rijke (2018)). There tends to be many causes, but one reason (which we explore in the paper) is due to the lack of statistical inferencing. The model itself lacks the possibilty of presenting a variety of lexical responses in the decoder, in a manner that a VAE could.
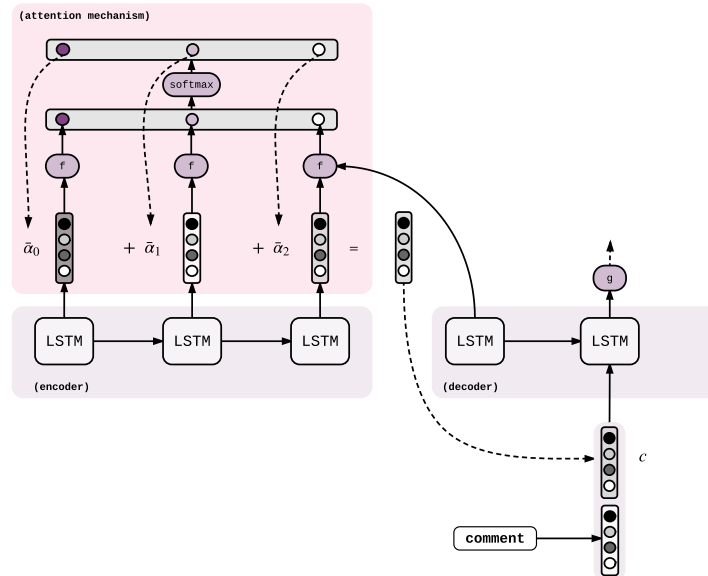
Attention Mechanism

Figure 2.7: An abstracted diagram of the attention mechanism, applied to a seq2seq model.

In addition to the context vector, the decoder can also retrieve non-contiguous relationships on the encoder values using attention (Bahdanau et al. (2014)). This mechanism looks at all of the inputs from the hidden states of the encoders so far. This allows the decoder network to "attend" to different parts of the source input at each step of the output generation. This would work in tandom with the context vector provided from the encoder, helping the model with long-range dependency problems.

Furthermore, this attention mechanism allows the model to learn what to attend to based on the input sequence and what it has so far, represented through a weighted combination of the two.

Attention mechanisms were found to perform particularly well in machine translation problems where languages which are relatively well aligned (such as English and German). The decoder is most likely able to choose to attend to the response

generation sequentially, such that the first output from the decoder can be generated based on the properties of the first input of the encoder and so on.

### 2.5.2 Variational Sequence To Sequence

Figure 2.8: The core structure of the VAE language model - words are represented as embedded vectors (Bowman et al. (2015)).

Instead of passing through the last hidden state drom the encoder to the decoder, it is possible to encode the context variable produced from the encoder in the form of gaussian parameters in a similar fashion to how VAEs are designed (Bowman et al. (2015)). The model is trained in a similar fashion to VAEs where the overall loss is composed of a reconstruction loss and a KL divergence (see Equation 2.7) but utilises additional mechanisms to force the language model to encode features into the gaussian parameters. We explain these features in further detail in Section 3.2. Although a neural language model based on this would be sufficient to generate responses, the responses would be generated from the same latent variable, restricting its variability in responses. (Zhao et al. (2017);Du et al. (2018))

2.5.3 Conditional Variational Sequence to Sequence

As established earlier, the VAD is an extension to the CVAE seq2seq model by Zhao et al. (2017), which itself is an adaptation of the original implementation by Bowman et al. (2015) for the purposes of discourse generation. It includes a plethora of additional components, including attention mechanisms to make discourse generation more viable.

We do not discuss this model in detail as it is rather similar to the VAD in terms of architectural design; Both models leverage a seq2seq base, variational latent variables, and a BOW loss mechanism. The primary difference is that the VAD samples the latent variable at each timestep of the decoder - the CVAE samples at the beginning of the decoder phase, and that the BOW loss for the VAD is called once to predict the whole decoder output.

Chapter 3

Model

In this chapter we discuss the variational autoregressive decoder (VAD) in detail. It builds on the knowledge found in the literature review, and the model itself is a combination of the components and models described above. To help fully understand the capabilities of the VAD, we subject it to a series of datasets to understand how it works and its performance against them. The VAD has been shown to encapsulate a distribution of the data within the latent parameters well compared to the CVAE Seq2Seq (Zhao et al. (2017)), and retains a high KL ratio respective to its overall loss function. This could potentially be used as the primary mechanism for encapsulating models, which can be sent to clients as a means of solving our case scenario.

3.1 Variational Autoregressive Decoders

Introduced by Du et al. (2018), VADs attempt to circumvent the sampling problem introduced from CVAEs by introducing multiple latent variables into the autoregressive Decoder. At different time-steps, this allows the decoder to produce a multimodal distribution of text sequences, allowing a greater variety of responses to be produced.

VADs use the vanilla seq2seq architecture as the base (Section 2.6) with variable-length queries $x = \{x_1, x_2, ..., x_n\}$, and $y = \{y_1, y_2, ..., y_n\}$ representing the query and response sequences respectively. Both the encoder and decoder utilises GRUs, with the encoder utilising a bidirectional GRU, and the decoder being unidirectional. For each timestep t, each GRU in the decoder network is encoded with hidden state h_t^d. The components of the model are broken down into their subsections below.

3.1.1 Encoder

$$
\begin{aligned}
\overrightarrow{h_t^e} &= \overrightarrow{GRU}(x_t, \overrightarrow{h_{t-1}^e}) \\
\overleftarrow{h_t^e} &= \overleftarrow{GRU}(x_t, \overleftarrow{h_{t+1}^e})
\end{aligned}
\tag{3.1}
$$

The encoder works in a similar fashion to the encoder described in the seq2seq network. The only changes here we see is that the encoder described in the implementation leverages bi-directionality. Similarly to the seq2seq and other recurrent
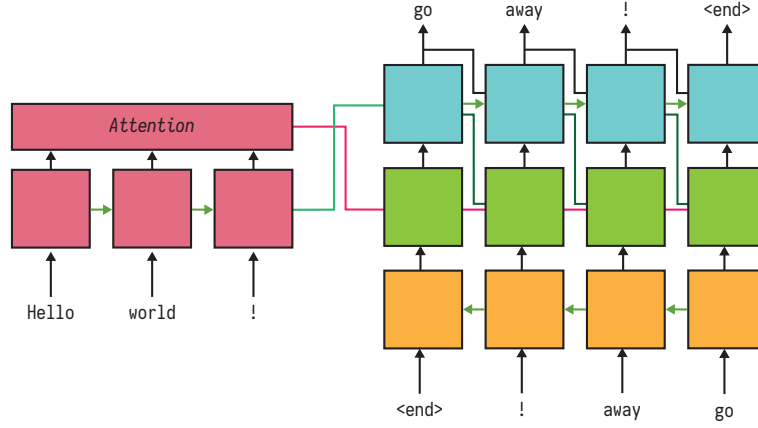
Figure 3.1: An abstracted model of the VAD model, where the encoder (pink) takes in the input sequence, orange represents the backwards model, green represents the latent, context generation, and the decoder (blue) shows the output sequence.

language models, the last hidden vector is passed through as the initial hidden vector for the decoder model.

3.1.2 Backwards and Attention

$$\alpha_{s,t} = f_{attention}([h^e_d, h^d_{t-1}])$$
$$\overleftarrow{h^d_t} = \overleftarrow{GRU}(y_{t+1}, \overleftarrow{h^d_{t+1}}) \qquad c_t = \sum_{s=1}^{m} \alpha_{s,t} h^e_s \qquad (3.2)$$

Figure 3.2: Equations for Backwards (left) and Attention (right).

The backwards RNN is used to feed additional contextual information for the inference model. During training, the backwards RNN takes as input the training outputs y and outputs hidden vectors in a reverse, sequential manner. The VAD leverages the attention mechanism in a similar manner, and takes in the input sequence, and is contextualised by the decoder at a given timestep.

3.1.3 Inference and Prior Models

$$[\mu^i, \sigma^i] = f_{infer}([\overrightarrow{h^d_{t-1}}, c_t, \overleftarrow{h^d_t}]) \qquad [\mu^p, \sigma^p] = f_{prior}([\overrightarrow{h^d_{t-1}}, c_t]) \qquad (3.3)$$
$$q_\theta(z_t|\boldsymbol{y}, \boldsymbol{x}) = \mathcal{N}(\mu^i, \sigma^i) \qquad p_\phi(z_t|\boldsymbol{y}_{<t}, \boldsymbol{x}) = \mathcal{N}(\mu^p, \sigma^p)$$

Figure 3.3: Equations for Inference (left) and Prior (right) models.

The inference model attempts to learn a conditional posterior distribution z using the output, the attention weights, and the decoder outputs. z is created using the gaussian reparameterisation trick described in Section 2.4.1.

This component stems away from Zhao et al. (2017), as $\boldsymbol{z}$ is now decomposed into sequential variables $\boldsymbol{z} = \{z_1, ..., z_t\}$, which are generated at each time step of the decoder phase. z_t is conditioned by the backwards hidden vector h_t^d. Du et al. (2018) suggests that this conditioning allows the latent variables to be guided for long-term generation.

The prior network is restricted to using observable variables during testing of the model to generate z_t using only observable data. It is designed in a similar fashion to the decoder network where the input variables are concatenated together. Note that the prior is no longer a standard gaussian; this particular component would learn to approximate the posterior distribution.

3.1.4   Decoder (Variational Autoregressive)

$$\overrightarrow{h_t^d} = \overrightarrow{GRU}([y_{t-1}, c_t, z_t], \overrightarrow{h_{t-1}^d})$$
$$p_\phi(y|\boldsymbol{y}_{<t}, \boldsymbol{z}_t, \boldsymbol{x}) = f_{output}([\overrightarrow{h_t^d}, c_t]) \tag{3.4}$$

The decoder then takes in the temporally conditioned z_t, alongside the previous output y_{t-1} and the context vector produced from the attention mechanism c_t to produce the output.

3.1.5   Auxillary Objective

One of the novel contributions of the model is the use of a temporally contextual auxiliary objective that uses the sampled latent vector z_t to predict the sequential bag of words (SBOW) of the response sequences $\boldsymbol{y}_{bow}(t+1, T)$. Let $f = MLP_b(z_t) \in \mathcal{R}^V$ where V be the vocabulary size, then the bag of words can be deduced using the following equation:

$$log[p_\xi(y_{bow(t+1,T)}|z_{t:T})] = log\frac{e^{f_{z_t}}}{\sum_j^V e^{f_j}} \tag{3.5}$$

The intuition is that this auxiliary objective would help to improve the KL divergence, as the latent variables would capture the information from a different perspective (Du et al. (2018)).

3.1.6   Learning Mechanism

VADs uses the ELBO (Equation 2.7)[1] and a weighted log likelihood loss of the auxillary objective, controlled by α. This loss is computed at each timestep, and is then

---

[1]As the KL divergence no longer compares the latent variable against a standard gaussian, we derive the equation to allow gradient propagation between two arbitary distribution. This can be seen in the Appendix (Section 7.1).

summed at the end of the sequence generation during the decoder stage. During training and evaluation, we sample both the inference and prior models to facilitate measuring the loss. The KL divergence allows the model to perform discourse generation using the prior mdoel, by forcing it to decode plausible sentences from every point in the latent space that has some reasonable odds under the prior (Bowman et al. (2015)).

$$\mathcal{L} = \sum_t [\mathcal{L}_{ELBO}(t) + \alpha \mathcal{L}_{AUX}(t)]$$
$$\mathcal{L} = \sum_t [(\mathcal{L}_{LL}(t) - \mathcal{L}_{KL}(t)) + \alpha \mathcal{L}_{AUX}(t)]$$

(3.6)

During implementation, it was found that the auxiliary loss weight α typically initially increases the KL divergence, but eventually produces a diminishing effect to both the auxiliary loss and the KL divergence. Attempts to contact the original authors of the VAD paper were to no avail, and is left for future work.

3.2 Optimisation Challenges

In addition to the fact that RNNs itself are difficult to train (see Section 2.1), it may occasionally be the case that with variational models, the sampled latent variable z is often ignored during training. Ideally, a model that encodes useful information in z will have a non-zero KL divergence term and in addition to a relatively small cross-entropy term.

Due to the sensitivity of the latent component, which is compounded by the sensitivity of the recurrent cells in the VAD, early builds of the VAD would often ignore z and go after the low hanging fruit during the learning process of the decoder. Straightforward implementations of our model failed to learn this behaviour; z would encorporate a lack of relevant information and is essentially ignored. That is not to say that the model does not learn - it would rather allocate more learning weight to the information from other inputs such as the attention mechanism and the encoder's hidden value.

The lack of learning in z results in a KL collapse such that the KL loss reaches zero, indicating a that the posterior distribution of z to "collapse" under the prior, effectively rendering z to be ignored, and thus the model is then effectively equivalent to an RNN language model. (Bowman et al. (2015)) We explain one of the mechanisms used to prevent KL collapse in the auxilary objective (see Section 3.1.5), but in this section, we discuss some additional approaches to overcome learning difficulties, primarily around z, but also for the model overall.

3.2.1 KL Annealing

A simple approach by Bowman et al. (2015), it involves adding a variable weight to the KL term in the cost function at training time. Initially, the weight is zero,

forcing the model to encode as much information into the latent variable as much as it can. The weight is then gradually increased, eventually reaching 1. At this stage the weighted cost function is equivalent to the ELBO loss. This could be thought of as annealing from a regular autoencoder towards a variational one.

That being said, Bowman et al. exposes a variety of additional hyperparameters including the step size, the total number of steps, the step rate (for instance, a linear progression or a logistic). For the sake of simplicity, we explore a linear approach with a fixed linear progression.

3.2.2 Word Dropout

Also proposed by Bowman et al. (2015), it involves weakening the decoder by removing some conditioning information during training. This is done by randomly replacing a fraction of the conditioned-on word tokens with a generic unknown token UNK, which forces the model to depend on the latent variable z for prediction. This works in a similar fashion to the standard dropout (Srivastava et al. (2014)) where connections are dropped between layers of a network, but is applied to the input data of a recurrent cell. We do not explore this technique further for the VAD as a case by Bowman et al. (2015) indicated it's distruptiveness when the drop rate is suffiently high.

3.2.3 Teacher Forcing

Teacher Forcing (Williams and Zipser (1989)) is a concept of using real target outputs as each next input in the decoder sequence as opposed to using the decoder's guess. This causes the model to converge faster, but may exhibit instability when the trained network is exploited. Typically, this can be controlled with a probability p such that the decoder sequence has a p chance of using teacher forcing during training.

Outputs can be observed when using teacher forced decoders that read with coherent grammar, but the generated response can wander far from the correct response; i.e. with some $p < 1$, some outputs from the decoder at some timestep would be correct, but the others would have a chance to completely change the semantics of the output. We initially set $p = 1$, but we set p to be explored during the hyperparameter search.

Chapter 4

Experimental Setup

We compare three neural language models: a baseline model, a VAE Seq2seq (see Section 2.5.2) and the VAD model. the baseline model represents a encoder-decoder seq2seq network (Serban et al. (2016)) that uses the same architecture base as the VAD but only minimises the standard cross entropy loss (i.e. the BOW loss and KL divergence computations are not computed). We use this approach instead of using a different implementation as it was discussed earlier in the paper that a KL collapse on variational models is effectively equivalent to a recurrent language model.

For all intents, the VAE Seq2Seq model is primarily designed for language modelling, whereas the VAD is designed for dialog response generation. With this prior knowledge it was expected that the VAE Seq2seq model would perform worse than the VAD model for all tasks.

The seq2seq model and the VAE Seq2seq models were chosen for comparision not necesssarily to compare state-of-the-art performance but to verify the correctness of the implementations against our VAD model. We do not implement the CVAE seq2seq model as it is not the focus of the ISO.

With the VAE seq2seq model, it is not necessarily designed with the intention of neural response generation, and that it was expected that the VAE seq2seq model would perform poorly in our scenario. However, to standardise the measurements for results, we pad the sequences such that it would still be useful for measuring the expected KL and reconstruction losses for a simple recurrent language model.

Weight Initialisation

Typically, we intend to have the weights asymmetricly generated (such that when we perform our gradient descent we would have random starting positions, which would enable us to have a better chance to find the minimum as opposed to having a fixed starting point).

Glorot and Bengio (2010) proposed to initialise the weights based on a gaussian $Var(w) = \mathcal{N}(0, \frac{2}{|n_{in}| + |n_{out}|})$ where w represents a layer in a network, and $|n_{in}|$ represents the number of neurons feeding into it, and similarly so for $|n_{out}|$. Note that I have used a gaussian distribution, but this could also be uniform. This method has shown a lot of success for sigmoidal and tanh based activation functions. We initialise all weights for all models with this approach.

4.1 Datasets

We subject the models to three different datasets: Amazon Reviews (He and McAuley (2016)), Penn TreeBank (Marcus et al. (2002)), and OpenSubtitles (Lison and Tiedemann (2016)). Each dataset serves a different purpose; we describe them in their respective sections.

4.1.1 Amazon Reviews

```
{
  'reviewerID': 'AA8JH8LD2H4P9',
  'asin': '7214047977',
  'reviewerName': 'Claudia J. Frier',
  'helpful': [3, 4],
  'reviewText': 'This fits my 7 inch kindle fire hd perfectly! I
     love it. It even has a slot for a stylus. The kindle is velcroed
     in so it\'s nice and secure. Very glad I bought this!',
  'overall': 5.0,
  'summary': 'love it',
  'unixReviewTime': 1354665600,
  'reviewTime': '12 5, 2012'
}
```

```
IDENTITY: ['7', '2', '1', '4', '0', '4', '7', '9', '7', '7',
     'rating_5.0', 'polarity_ -0.1']

['<sos>', 'love', 'it', '<sor>']
['this', 'fits', 'my', '7', 'inch', 'kindle', 'fire', 'hd',
     'perfectly', '!', 'i', 'love', 'it', '.']
['it', 'even', 'has', 'a', 'slot', 'for', 'a', 'stylus', '.']
['the', 'kindle', 'is', 'velcroed', 'in', 'so', 'it', "'", 's',
     'nice', 'and', 'secure', '.']
['very', 'glad', 'i', 'bought', 'this', '!', '<eos>']
```

Figure 4.1: A sample review and the augmented data. Note that for each sequence from line 3 onwards has the identity sequence contatenated before it.

The Amazon products reviews (He and McAuley (2016)) comprises a set of customer submitted reviews of products on the popular consumer commerce website Amazon.com, spanning May 1996 - July 2014. For the purposes of this ISO, we select the electronic products dataset, spanning 1.6m reviews, but due to computational constraints, we work on a subset of this dataset as reviews would also need to be broken down into sequences of sentences. This dataset is used as the primary indicator of the performance of the models in their capability for text generation, and the capabilities of encompassing lexical variety, as described in the original objective of the ISO. The objective of the models were to produce the next sentence from the previous sentence.

Pre-processing includes (1) sentenisation and tokenisation using SpaCy (Honnibal and Montani (2017)); (2) concatenation of sequence conditioning on the input se-

quences; (3) keeping the top 10K frequent word types af the vocabulary; (4) removing sequences where the "¡unk¿" tags in the sequences represent a ratio of greater than 0.1; (5), and finally removing sequences that are greater than a length of 60. Each input sentence is conditioned by the Item ID, review rating (of a score from 1-5), and a polarity value consisting of -1 to $+1$ with an increment of 0.1.

GloVe Word embeddings (Pennington et al. (2014)) are used to represent tokens in the dataset, and is also used to determine words to remove. The polarity value is computed p with a rounded ratio $p = round(r_+/r_-, 2)$. The polarity word vector is a sum of the polarity value p and the word vector corresponding to the word "polarity". This works in a similar fashion for the rating word embedding. Item IDs are split character wise.

4.1.2 OpenSubtitles

```
1 "<sos> ya ! <eos>" -> "<sos> write to your <unk> and hurry that up .
    <eos>"
2 "<sos> your paycheck ? <eos>" -> "<sos> back off tucker , you don '
    t sketch regulations . <eos>"
```

Figure 4.2: A sample review and the augmented data. Note that for each sequence from line 3 onwards has the identity sequence contatenated before it.

The OpenSubtitles dataset (Lison and Tiedemann (2016)) is used as a measurement in a similar fashion to the amazon dataset but does not involve any conditioning on the input variables. They are setup in a similar fashion to the amazon dataset, described above. GloVe is also used. We use this to measure the effectiveness of the models on discourse generation without the use of additional contextual information. The goal for the models was to predict the next subtitle line, given a previous subtitle line.

Pre-processing works in a similar fashion to the amazon dataset. It includes (1) sentenisation and tokenisation using SpaCy (Honnibal and Montani (2017)); (2) concatenation of sequence conditioning on the input sequences; (3) keeping the top 10K frequent word types af the vocabulary; (4) removing sequences where the "¡unk¿" tags in the sequences represent a ratio of greater than 0.1; (5), and finally removing sequences that are greater than a length of 60.

4.1.3 Penn TreeBank

For quick configuration with the model, we compute initial results with the Penn Tree-Bank Dataset. (Marcus et al. (2002)) This dataset is typically used for verification of model performance for generative models; the dataset consists of sequences which are the same for the input and output. (See 4.3) For the purposes of the ISO, we used the exemplar dataset which consists of 43K sequences. Pre-processing includes (1) tokenisation using NLTK; (2) trimming sequences that are greater than a length of 60; (3) replacing tokens with ¡unk¿ where the token does not exist in

```
1 [['<sos>', 'the', 'n', 'stock', 'specialist', 'firms', 'on', 'the',
      'big', 'board', 'floor', 'the', 'buyers', 'and', 'sellers', 'of',
      'last', 'resort', 'who', 'were', 'criticized', 'after', 'the',
      'n', 'crash', 'once', 'again', 'could', "n't", 'handle', 'the',
      'selling', 'pressure'], ['the', 'n', 'stock', 'specialist',
      'firms', 'on', 'the', 'big', 'board', 'floor', 'the', 'buyers',
      'and', 'sellers', 'of', 'last', 'resort', 'who', 'were',
      'criticized', 'after', 'the', 'n', 'crash', 'once', 'again',
      'could', "n't", 'handle', 'the', 'selling', 'pressure', '<eos>']]
2 [['<sos>', 'big', 'investment', 'banks', 'refused', 'to', 'step',
      'up', 'to', 'the', 'plate', 'to', 'support', 'the',
      'beleaguered', 'floor', 'traders', 'by', 'buying', 'big',
      'blocks', 'of', 'stock', 'traders', 'say'], ['big', 'investment',
      'banks', 'refused', 'to', 'step', 'up', 'to', 'the', 'plate',
      'to', 'support', 'the', 'beleaguered', 'floor', 'traders', 'by',
      'buying', 'big', 'blocks', 'of', 'stock', 'traders', 'say',
      '<eos>']]
```

Figure 4.3: A sample review and the augmented data. Note that for each sequence from line 3 onwards has the identity sequence contatenated before it.

the vocabulary. For this particular dataset, we do not use any pre-trained word-embeddings; as the purposes of this dataset does not necessarily benefit from its addition. The purpose of this dataset is to measure the models reconstructive capabilities as a regular autoencoder. Models would use this dataset to repeat the input as the output sequence.

4.2 Hyperparameter Setup

For the sake of comparisons, all models leverage a bidirectional encoder. For variational models we utilise a linear KL annealing with a step size of $2.5 \cdot 10^{-5}$ with a threshold of 2000 steps, both are which are hyperparameterised. Dimension sizes for the hidden sizes are 512, with latent sizes being 400, and are not hyperparameterised. The cross entropy loss, KL loss, and BOW losses are normalised by the sequence length.

All models use the Adam optimiser (Kingma and Ba (2014)) with a base learning rate of 0.001; subject to hyperparameterisation. As we measure the performance of three fundamentally different neural language architectures, we performed hyperparameter tuning for each of the models independently.

Datasets were split in a 80:15:5 ratio, for training, validation, and test data respectively. Models were subject to roughly 15 hyperparameter tuning rounds (dependent on the confidence of the optimisation model) using Bayesian Optimisation (GPyOpt authors (2016)) where the objective is to minimise the loss. Early stopping is initialised such that the training will end when the mean loss of the validation batch is reached below an overall loss of 1.0.

4.3 Measuring Performance

4.3.1 BLEU and ROUGE

It is common in the popular cases to use subject human guided hand-picked results to emperical measurements but due to a variety of constraints for the ISO this would not be feasible. Alternatively, we subject the models generated responses to the BLEU (Papineni et al. (2001)) and ROUGE (Lin (2004)) emperical scores at each epoch to quantify the quality of our responses. These measurements are used alongside the Loss, and human analysis of the results.

BLEU (Papineni et al. (2001)) is originally defined to measure n-gram precision p_n of two text based sequences, by summing the n-gram matches for every hypothesis sentence S in the test corpus C. ROUGE is used in a similar fashion for recall. Both BLEU and ROUGE have the option to compute the n-gram match for several sizes of n-grams, but for simplicity we reduce it to uni-grams.

$$p_n = \frac{\sum_{S \in C} \sum_{ngram \in S} Count_{clip}(ngram)}{\sum_{S \in C} \sum_{ngram \in S} Count(ngram)} \qquad p_n = \frac{\sum_{S \in C} \sum_{ngram \in S} Count_{matched}(ngram)}{\sum_{S \in C} \sum_{ngram \in S} Count(ngram)}$$

$$(4.1)$$

Figure 4.4: Equations for BLEU (left) and ROUGE (right).

Although both equations look very similar, BLEU introduces a brevity penalty term where it would penalise results that are shorter than the general length of the test corpus.

$$f1_n = 2 \frac{BLEU_n \cdot ROUGE_n}{BLEU_n + ROUGE_n}$$

$$(4.2)$$

Figure 4.5: The harmonic mean of BLEU and ROUGE scores.

To understand an overall picture of the model performance, we also measure the harmonic mean of BLEU and ROUGE, by considering the F1 performance. We consider that the larger the F1 performance, the more faithful the model is in replicating the sequences of the output data given its conditioning.

Chapter 5

Experimental Results

5.1 NLL Loss

As we compare loss values for variational and non-variational models, we split the loss values into their constituents and analyse from there to provide a better understanding of the performance of the models. The measurements below are taken from the average loss over the training batches for each epoch.


Figure 5.1: Reconstruction losses of the three models across the three datasets; lower loss is better.

The results shown are not necessarily surprising. As there are less weights to learn, it was expected as the vanilla seq2seq model to converge faster than the other models. Conversely, it was expected for the VAD model to converge the slowest considering the number of weights involved. Additionally, varational models are expected to perform worse than discrete models as the sampling mechanism will produce some variation in the responses.

Interestingly, for the amazon dataset, we observe that the bowman model converges very quickly, but later on we view that the bowman model has overfitted on the training sequence.

Figure 5.2: KL losses of the three models across the three datasets; higher is better.

5.2 KL Loss

We measure the KL loss to determine the rate of which the amount of information is stored in the latent variable. Typically, we expect the models to produce some non-zero KL response, but we have noticed examples of KL collapse with the bowman approach. With the VAD, we observe larger KL values relative to the bowman model, suggesting that the auxiliary function works as intended for all types of datasets.

In a similar fashion, the bowman model is expected to perform weaker than the VAD approach as the objective of the latent vector is to encompass properties of the data from a global sequence.

5.3 KL Ratio

Figure 5.3: KL ratios of the three models across the three datasets; higher is better.

The intent of measuring the KL Ratio is such that we would be able to see a good picture of how much information was encompassed in the latent variable, and thus retaining more characteristics of the variational autoencoder as opposed to a regular autoencoder. Interestingly, it was easier for the VAD to initially encode more information into the latent variable for the Penn dataset.

5.4 BLEU and Rouge

Note that from this section onwards, we will be looking at the performance against the test splits of the datasets.
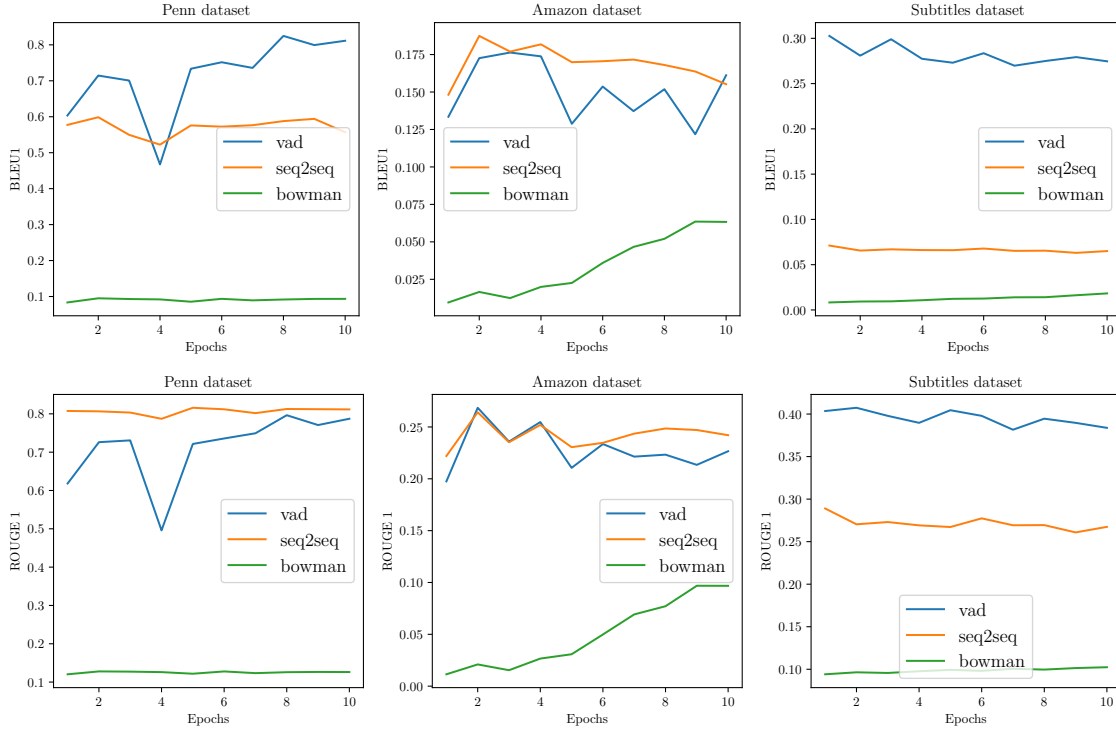

Figure 5.4: BLEU$_1$ (top) and ROUGE$_1$ (bottom) scores across the three models; higher is better.

With the BLEU scores, Surprisingly, for the Penn Dataset, the VAD is considerably stronger than the other models by a large margin, suggesting that it performs very strongly as an autoencoder. With the Amazon dataset,
We expected that the Bowman model would perform the weakest considering that the loss mechanism does not consider output sequences of arbitrary sizes.
Interestingly, we observe that the bowman model overfits heavily, with a low BLEU and ROUGE score, in comparison to the reconstruction losses observed on the training data. This is most likely caused by the considerably smaller dataset than what was used in the reference paper, where data sizes ranged towards the 80 million input and output sequences. With that being said, the focus of the paper again revolves around the VAD, which for our purposes have shown to perform well.

5.5 F1

5.6 Output Variance

The motive of the responses is to The responses produced here should com

Figure 5.5: $F1_{\text{1-gram}}$ scores across the three models; higher is better.

5.7 Sampling Examples

In additional to quantitiative measurements, we also generate example resonses from the models to determine how they fare against the reference response in each of the datasets.

As some of our models are based on bayesian inference, we sample five outputs and hand pick the most optimal responses out of each of the relevant models.

Chapter 6

Future Work

- Generated responses may not necessarily be indicative of the data it is modelled against.

- Although there are mechanisms to control the vanishing KL problem, it is yet to be seen how universally applicable it would be.

- The amount of hyperparameter tuning is considerably larger than other neural models due to the inclusion of additional mechanisms for controlling the vanishing KL.

- further testing of the additional components introduced by the VAE should take place to measure their influence to the KL loss. - Statistical inference should be used to understand the direction of the KL divergence further. -
One of the caveats of a research is typically that research direction is open ended. - talk about how you could potentially increase the KL weight by merging the SBOW with the reconstruction loss
- use nn to classify performance of outputs to detect realisticness of data - apply to gan to improve model performance - more measurements - BLEU/ROUGE on the most output instead of an average of the whole sequence but requires extensive memory usage and would take a considerably longer running time due to technical implementation. - longer Hyperparameter optimisation - augment amazon reviews s.t. the item ID is the mean of the item weights. - explore completely generative sequences end to end.

Chapter 7

Conclusion

In conclusion, we explored a neural language model that could potentially be used for the purposes of encompassing the lexical variety of the original dataset. By understanding the limitations of the model and how to augment the components in a manner that allow them to work in unison, we conclude that through quantitiative and qualitative experimental results, alongside other references, the VAD can be sensibly considered for the task at hand. In future works, the model could be extended such that it would produce stronger and more lexically coherent discourses in a similar fashion as the dataset. One avenue for exploration could be related to adversarial approaches.

Bibliography

Bahdanau, D., Cho, K., and Bengio, Y. (2014). Neural Machine Translation by Jointly Learning to Align and Translate. *arXiv:1409.0473 [cs, stat]*. arXiv: 1409.0473. pages 11

Bengio, Y., Ducharme, R., and Vincent, P. (2001). A Neural Probabilistic Language Model. In Leen, T. K., Dietterich, T. G., and Tresp, V., editors, *Advances in Neural Information Processing Systems 13*, pages 932–938. MIT Press. pages 5

Bowman, S. R., Vilnis, L., Vinyals, O., Dai, A. M., Jozefowicz, R., and Bengio, S. (2015). Generating Sentences from a Continuous Space. *arXiv:1511.06349 [cs]*. arXiv: 1511.06349. pages 12, 16, 17

Cho, K., van Merrienboer, B., Bahdanau, D., and Bengio, Y. (2014). On the Properties of Neural Machine Translation: Encoder-Decoder Approaches. *arXiv:1409.1259 [cs, stat]*. arXiv: 1409.1259. pages 7

Chung, J., Gulcehre, C., Cho, K., and Bengio, Y. (2014). Empirical Evaluation of Gated Recurrent Neural Networks on Sequence Modeling. *arXiv:1412.3555 [cs]*. arXiv: 1412.3555. pages 7

Du, J., Li, W., He, Y., Xu, R., Bing, L., and Wang, X. (2018). Variational Autoregressive Decoder for Neural Response Generation. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing*, pages 3154–3163, Brussels, Belgium. Association for Computational Linguistics. pages 2, 12, 13, 15

Dyer, C. (2017). Conditional Language Modelling. original-date: 2017-02-06T11:32:46Z. pages 4

E. Rumelhart, D., H. Hinton, G., and RJ, W. (1986). Learning Internal Representations by Error Propagation. *Parallel Distributed Processing: Explorations in the Microstructure of Cognition*, 1. pages 7

Fu*, H., Li*, C., Liu, X., Gao, J., Celikyilmaz, A., and Carin, L. (2019). Cyclical Annealing Schedule: A Simple Approach to Mitigating KL Vanishing. pages 8

Glorot, X. and Bengio, Y. (2010). Understanding the difficulty of training deep feedforward neural networks. In *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, pages 249–256. pages 18

GPyOpt authors (2016). GPyOpt: A Bayesian Optimization framework in python. pages 21

He, R. and McAuley, J. (2016). Ups and Downs: Modeling the Visual Evolution of Fashion Trends with One-Class Collaborative Filtering. In *Proceedings of the 25th International Conference on World Wide Web - WWW '16*, pages 507–517, Montréal, Québec, Canada. ACM Press. pages 2, 19

Hochreiter, S. and Schmidhuber, J. (1997). Long Short-Term Memory. *Neural Computation*, 9(8):1735–1780. pages 6

Honnibal, M. and Montani, I. (2017). spaCy 2: Natural language understanding with Bloom embeddings, convolutional neural networks and incremental parsing. pages 19, 20

Jiang, S. and de Rijke, M. (2018). Why are Sequence-to-Sequence Models So Dull? Understanding the Low-Diversity Problem of Chatbots. In *Proceedings of the 2018 EMNLP Workshop SCAI: The 2nd International Workshop on Search-Oriented Conversational AI*, pages 81–86, Brussels, Belgium. Association for Computational Linguistics. pages 11

Jurafsky, D. and Martin, J. (2019). *Speech and Language Processing*. 3 edition. pages 5

Kannan, A., Young, P., Ramavajjala, V., Kurach, K., Ravi, S., Kaufmann, T., Tomkins, A., Miklos, B., Corrado, G., Lukacs, L., and Ganea, M. (2016). Smart Reply: Automated Response Suggestion for Email. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining - KDD '16*, pages 955–964, San Francisco, California, USA. ACM Press. pages 4

Karpathy, A. (2015). The Unreasonable Effectiveness of Recurrent Neural Networks. pages 5

Kingma, D. P. and Ba, J. (2014). Adam: A Method for Stochastic Optimization. pages 21

Kingma, D. P. and Welling, M. (2013). Auto-Encoding Variational Bayes. *arXiv:1312.6114 [cs, stat]*. arXiv: 1312.6114. pages 2, 8, 9

Kovalenko, B. (2017a). *Controllable text generation*. PhD thesis, Stanford University. pages 3

Kovalenko, B. (2017b). Controllable text generation. pages 3

Le, J. (2018). Recurrent Neural Networks: The Powerhouse of Language Modeling. pages 4

Lin, C.-Y. (2004). ROUGE: A Package for Automatic Evaluation of Summaries. pages 74–81. pages 22

Lison, P. and Tiedemann, J. (2016). OpenSubtitles2016: Extracting Large Parallel Corpora from Movie and TV Subtitles. page 7. pages 2, 19, 20

Marcus, M., Ann Marcinkiewicz, M., and Santorini, B. (2002). Building a Large Annotated Corpus of English: The Penn Treebank. *Computational Linguistics*, 19:313–330. pages 19, 20

Nallapati, R., Zhai, F., and Zhou, B. (2016). SummaRuNNer - A Recurrent Neural Network based Sequence Model for Extractive Summarization of Documents. pages 4

Papineni, K., Roukos, S., Ward, T., and Zhu, W.-J. (2001). BLEU: a method for automatic evaluation of machine translation. In *Proceedings of the 40th Annual Meeting on Association for Computational Linguistics - ACL '02*, page 311, Philadelphia, Pennsylvania. Association for Computational Linguistics. pages 22

Pennington, J., Socher, R., and Manning, C. D. (2014). GloVe: Global Vectors for Word Representation. In *Empirical Methods in Natural Language Processing (EMNLP)*, pages 1532–1543. pages 20

Raiko, T., Berglund, M., Alain, G., and Dinh, L. (2014). Techniques for Learning Binary Stochastic Feedforward Neural Networks. *arXiv:1406.2989 [cs, stat]*. arXiv: 1406.2989. pages 8

Serban, I. V., Sordoni, A., Lowe, R., Charlin, L., Pineau, J., Courville, A., and Bengio, Y. (2016). A Hierarchical Latent Variable Encoder-Decoder Model for Generating Dialogues. *arXiv:1605.06069 [cs]*. arXiv: 1605.06069. pages 11, 18

Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2014). Dropout: A Simple Way to Prevent Neural Networks from Overfitting. *Journal of Machine Learning Research*, 15:1929–1958. pages 17

Sutskever, I., Vinyals, O., and Le, Q. V. (2014). Sequence to Sequence Learning with Neural Networks. *arXiv:1409.3215 [cs]*. arXiv: 1409.3215. pages 4, 10

Williams, R. J. and Zipser, D. (1989). A Learning Algorithm for Continually Running Fully Recurrent Neural Networks. *Neural Computation*, 1(2):270–280. pages 17

Zhao, T., Zhao, R., and Eskenazi, M. (2017). Learning Discourse-level Diversity for Neural Dialog Models using Conditional Variational Autoencoders. *arXiv:1703.10960 [cs]*. arXiv: 1703.10960. pages 11, 12, 13, 15

Appendix

7.1   KL Divergence Derivation

Let $\mu_1, \sigma_1 \to \mathcal{N}(\mu_1, \sigma_1)$ be our first distribution, and $\mu_2, \sigma_2 \to \mathcal{N}(\mu_2, \sigma_2)$ be our second distribution. By deriving this, we would be able to calculate a derivative friendly loss function for our models.

$$
\begin{aligned}
KL &= \int \left[\log(p(x)) - \log(q(x)) \right] \; p(x) \; dx \\
&= \int \left[\frac{1}{2} \log \frac{|\Sigma_2|}{|\Sigma_1|} - \frac{1}{2}(x - \mu_1)^T \Sigma_1^{-1}(x - \mu_1) + \frac{1}{2}(x - \mu_2)^T \Sigma_2^{-1}(x - \mu_2) \right] \times p(x) dx \\
&= \frac{1}{2} \log \frac{|\Sigma_2|}{|\Sigma_1|} - \frac{1}{2}\mathrm{tr} \left\{ E[(x - \mu_1)(x - \mu_1)^T] \, \Sigma_1^{-1} \right\} + \frac{1}{2} E[(x - \mu_2)^T \Sigma_2^{-1}(x - \mu_2)] \\
&= \frac{1}{2} \log \frac{|\Sigma_2|}{|\Sigma_1|} - \frac{1}{2}\mathrm{tr} \left\{ I_d \right\} + \frac{1}{2}(\mu_1 - \mu_2)^T \Sigma_2^{-1}(\mu_1 - \mu_2) + \frac{1}{2}\mathrm{tr}\{\Sigma_2^{-1}\Sigma_1\} \\
&= \frac{1}{2} \left[\log \frac{|\Sigma_2|}{|\Sigma_1|} - d + \mathrm{tr}\{\Sigma_2^{-1}\Sigma_1\} + (\mu_2 - \mu_1)^T \Sigma_2^{-1}(\mu_2 - \mu_1) \right].
\end{aligned}
$$

$$(7.1)$$